

TCP Retransmission and Flow Control (Part II)

Presenter: Assist.Prof.Dr. Fatih ABUT

What TCP does for you?

- Stream-based in-order delivery
 - Segments are ordered according to sequence numbers
 - Only consecutive bytes are delivered
- Reliability
 - Missing segments are detected (ACK is missing) and retransmitted
- Flow control
 - Receiver is protected against overload (window based)
- Congestion control
 - Network is protected against overload (window based)
 - Protocol tries to fill available capacity
- Connection handling
 - Explicit establishment + teardown
- Full-duplex communication
 - e.g., an ACK can be a data segment at the same time (piggybacking)

TCP Retransmission Strategy

- TCP relies on positive acknowledgements
 - Retransmission on timeout
- Timer associated with each segment as it is sent
- If timer expires before acknowledgement, sender must retransmit
- TCP uses an *adaptive retransmission algorithm* because internet delays are so variable
- *Round trip time* of each connection is recomputed every time an acknowledgment arrives
- Timeout value is adjusted accordingly

Retransmit Timeout Mechanism (1)

- RTO timer value difficult to determine:
 - too high \Rightarrow bad in case of msg-loss!
 - too short \Rightarrow risk of false alarms!
 - General consensus: too short is worse than too long; use conservative estimate
- Calculation: measure RTT (Seg# ... ACK#)
- Original suggestion in RFC 793 (1981)
 - Exponentially Weighed Moving Average (EWMA)
 - $SRTT_{new} = \alpha * SRTT_{old} + (1-\alpha) * RTT_{current}$
 - $RTO = \min(UBOUND, \max(LBOUND, \beta * SRTT))$

SRTT: Smoothed Round Trip Time

α : Smoothing factor (typically 0.8-0.9)

β : Variance factor (typically 2)

Retransmit Timeout Mechanism (2)

- Depending on variation, this RTO may be too small or too large; thus, final algorithm includes deviation
- Key observation:
 - Original smoothed RTT can't keep up with wide/rapid variations in RTT
- Jacobson/Karel's Retransmission Timeout (1988)


- $SRTT_{new} = \alpha * SRTT_{old} + (1-\alpha) * RTT_{current}$
- $SDEV_{new} = \beta * SDEV_{old} + (1 - \beta) * |\text{abs}(RTT_{current} - SRTT_{new})|$
- $RTO_{new} = SRTT_{new} + \gamma * SDEV_{new}$

SDEV: Smoothed deviation

β : Smoothing factor for standard deviation (typically 0.8-0.9)


γ : Adjustment factor (typically 4)

RTO calculation



- Problem: retransmission ambiguity
 - Segment #1 sent, no ACK received → segment #1 retransmitted
 - Incoming ACK #2: cannot distinguish whether original or retransmitted segment #1 was ACKed
 - Thus, cannot reliably calculate RTO!
- Solution [Karn/Partridge]: ignore RTT values from retransmits
 - Problem: RTT calculation especially important when loss occurs; sampling theorem suggests that RTT samples should be taken more often
- Solution: Timestamps option
 - Sender writes current time into packet header (option)
 - Receiver reflects value
 - At sender, when ACK arrives, $\text{RTT} = (\text{current time}) - (\text{value carried in option})$


Sliding Window Management



- Receiver “grants” credit (receiver window, rwnd)
 - sender restricts sent data with window
- Receiver buffer not specified
 - i.e. receiver may buffer reordered segments (i.e. with gaps)

Small Packet Problem

- Sending data in very small segments
 - Syndrome created by the Sender
 - Sending application program creates data slowly (e.g. 1 byte at a time)
 - Wait and collect data to send in a larger block
 - How long should the sending TCP wait?
 - Solution: Nagle's algorithm
 - Nagle's algorithm takes into account (1) **the speed of the application program that creates the data, and (2) the speed of the network that transports the data**



Nagle's Algorithm (1)

- If there is data to send but the window is open less than MSS, then we may want to wait some amount of time before sending the available data
 - If we wait too long, then we hurt interactive applications
 - If we don't wait long enough, then we risk sending a bunch of tiny packets and falling into the silly window syndrome
- The solution is to introduce a timer and to transmit when the timer expires

Nagle's Algorithm (2)

- We could use a clock-based timer, for example one that fires every 100 ms
- Nagle introduced an elegant self-clocking solution
- Key Idea
 - As long as TCP has any data in flight, the sender will eventually receive an ACK
 - This ACK can be treated like a timer firing, triggering the transmission of more data

Nagle's Algorithm (3)

When the application produces data to send

 if both the available data and the window \geq MSS

 send a full segment

 else /* window < MSS */

 if there is unACKed data in flight


 buffer the new data until an ACK arrives

 else

 send all the new data now

Silly Window Syndrome

- Syndrome created by the Receiver
 - Receiving application program consumes data slowly (e.g. 1 byte at a time)
 - The receiving TCP announces a window size of 1 byte. The sending TCP sends only 1 byte.
 - Solution: Sending an ACK but announcing a window size of zero until there is enough space to accommodate a segment of max. size or until half of the buffer is empty



Fast Retransmit (1)

- Coarse timeouts remained a problem, and **Fast retransmit** was added with TCP Tahoe.
- Since the receiver responds every time a packet arrives, this implies the sender will see duplicate ACKs.
- Basic Idea: *use **duplicate ACKs** to signal lost packet.*

Fast Retransmit

Upon receipt of **three** duplicate ACKs, the TCP Sender retransmits the lost packet.

Fast Retransmit (2)

- Generally, **fast retransmit** eliminates about half the coarse-grain timeouts.
- This yields roughly a 20% improvement in throughput.

Fast Retransmit

Based on three duplicate ACKs

